



Shinystan module structure

Outline

- Shinytan 3.0
 - Backend changes
 - Modules, why?
 - Report generation
- On workflow
 - Where does shinytan fit in?
 - Adding appendices

Why?

- Make backend more stable / sustainable
- Easier to add / remove features
 - No naming issues due to different environments
- Base everything on bayesplot
- Add report option

global.r

server.r

ui.r

global.r

server.r

ui.r

Diagnose

Estimate

Home

More

Utils

global.r

server.r

ui.r

Diagnose

Estimate

Home

More

Utils

Lowest level
modules

Lowest level
modules

Lowest level
modules

Lowest level
modules

Lowest level
modules

How do modules work?

- Write modules which have UI and Server part
- `source("MODULES/DIAGNOSE/acceptance.r", local = TRUE)`

UI

```
acceptanceUI <- function(id){  
  ns <- NS(id)  
  ....  
}
```

NESTED:

```
...  
fluidRow(  
  align = "right",  
  plotOptionsUI(ns("options"))  
)  
...
```


SERVER

```
acceptance <- function(input, output, session){  
  ...  
}
```

NESTED:



```
...  
visualOptions <- callModule(plotOptions, "options")  
...
```

What can you assume?

- Shinystan object
 - shinystan::as.shinystan()

model_name	posterior_sample	n_chain	user_model_info
param_names	summary	n_iter	model_code
param_dims	sampler_params	n_warmup	misc

Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾



- 65 of 4000 iterations ended with a divergence (1.6%).
- 1 parameters have an effective sample size less than 10% of the total sample size.

DIAGNOSE

ESTIMATE

MORE

Close

SHINYSTAN

DIAGNOSE

ESTIMATE

MORE

Plots

Stats

Report

NUTS/HMC

Divergent Scatter

Parallel Coordinates

Pairs

Divergence Information

Energy Information

Treedepth Information

Step Size Information

Acceptance Information

MCMC

Autocorrelation

Trace Plots

Rank Plots

\hat{R} , n_{eff} , se_{mean}

Parameter

tau log-posterior

Transform X

identity


Transform Y

identity

All chains

0



Show/Hide Options



Show/Hide Caption

Include in report?

Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾

Plots

Stats

Report

NUTS/HMC

Divergent Scatter

Parallel Coordinates

Pairs

Divergence Information

Energy Information

Treedepth Information

Step Size Information

Acceptance Information

MCMC

Autocorrelation

Trace Plots

Rank Plots

\hat{R} , n_{eff} , se_{mean}

Parameter

tau log-posterior

All chains

0

☒ Show/Hide Options

Select Theme


classic ▾

Select Colors


viridisA ▾

Select Divergent Color

red ▾





☐ Show/Hide Caption



☐ Include in report?

Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾

Plots

Stats

Report

NUTS/HMC

Divergent Scatter

Parallel Coordinates

Pairs

Divergence Information

Energy Information

Tredepth Information

Step Size Information

Acceptance Information

MCMC

Autocorrelation

Trace Plots


Rank Plots

\hat{R} , n_{eff} , se_{mean}

All chains

0

☐ Show/Hide Options



☒ Show/Hide Caption

These are plots of the *acceptance statistic* (top leftpanel), the *log-posterior* (top right panel), and, the *acceptance statistic* (x-axis bottom panel) against the *log-posterior* (y-axis bottom panel) for all chains. The vertical lines indicate the mean (solid line) and median (dashed line). A bad plot would show a relationship between the acceptance statistic and the log-posterior. This might be indicative of poor exploration of parts of the posterior which might be mitigated by reparametrization or adaptation of the step size. If many proposals are rejected the integrator step size might be too large and the posterior might not be fully explored. If the acceptance rate is very high this might be indicative of inefficient sampling. The target Metropolis acceptance rate can be set with the *adapt_delta* control option. For more information see https://mc-stan.org/docs/2_19/reference-manual/hmc-algorithm-parameters.html.

☒ Include in report?

Close

SHINY STAN  DIAGNOSE  ESTIMATE MORE ▾

Plots

Stats

Report

NUTS/HMC

All NUTS/HMC stats

MCMC

\hat{R}_i , n_{eff} , se_{mean}

Statistic

☒ Mean ☐ SD ☐ Max ☐ Min

Warmup

☒ Omit ☐ Include

Decimals

2

	accept_stat	stepsize	treedepth	n_leapfrog	divergent	energy
All chains	0.83	0.24	3.64	17.44	0.02	24.00
chain1	0.74	0.34	3.12	10.89	0.02	24.37
chain2	0.93	0.12	4.36	27.25	0.00	23.40
chain3	0.90	0.16	3.97	20.90	0.02	23.69
chain4	0.73	0.35	3.08	10.73	0.02	24.54


Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾

Plots Stats Report



ShinyStan Diagnostics Report



This is a plot of MCMC draws of τ (x-axis) against log-posterior (y-axis). The red colored draws represent, if present, divergent transitions. Divergent transitions can indicate problems for the validity of the results. A good plot would show no divergent transitions. A bad plot would show divergent transitions in a systematic pattern. For more information see <https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>.

Close

Plots

Stats

Report

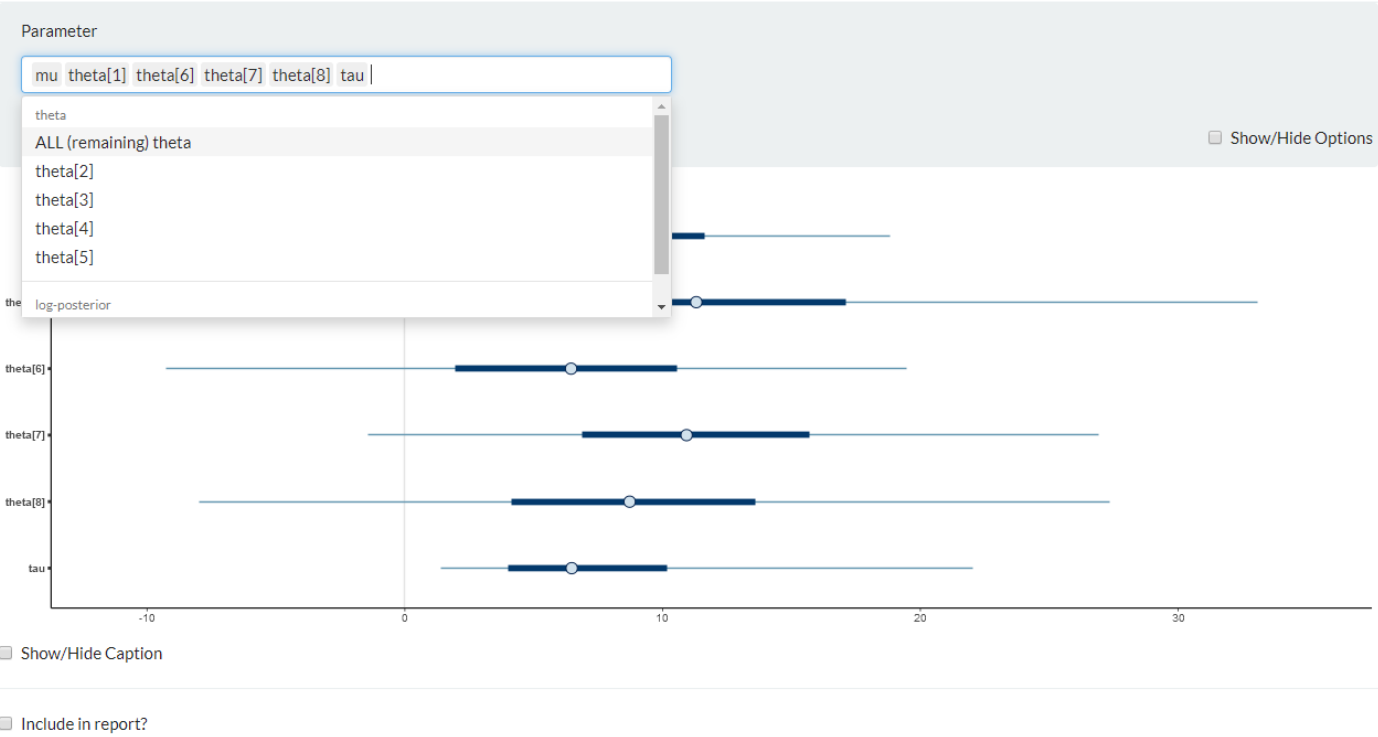
Intervals

Areas

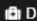

Scatter

Density

Histogram



Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾

Plots

Stats

Report

Intervals

Areas

Scatter

Density

Histogram

Parameter

mu theta[1] theta[2] theta[3] theta[4] theta[5] theta[6] theta[7] theta[8] tau

☒ Show/Hide Options

Select Theme

bayesplot default ▾

Select Colors

blue ▾

Point Estimate

☒ Median ☐ Mean ☐ None

Plot Type


☐ Areas ☒ Ridges

Posterior Interval (inner)

50% 95%

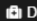

Posterior Interval (outer)

50% 95%



☐ Show/Hide Caption

Close

SHINYSTAN  DIAGNOSE  ESTIMATE MORE ▾

Plots

Stats

Report

Posterior Summary Statistics

Generate LaTeX table

Parameter

mu theta[1] theta[2] theta[3] theta[4] theta[5] theta[6] theta[7] theta[8] tau

Caption

Include a caption.

Decimals

1

Columns

- ☒ Posterior mean
- ☐ Monte Carlo error
- ☒ Posterior standard deviation
- ☒ Quantile: 2.5%
- ☐ Quantile: 25%
- ☒ Quantile: 50%
- ☐ Quantile: 75%
- ☒ Quantile: 97.5%
- ☐ Effective sample size
- ☒ Rhat

Packages

- ☐ Booktabs
- ☐ Longtable

```
% latex table generated in R 3.6.1 by xtable 1.8-3 package
% Mon Sep 16 12:19:37 2019
\begin{table}[ht]
\centering
\begin{tabular}{rllllll}
\hline
& mean & sd & 2.5\% & 50\% & 97.5\% & Rhat \\\
\hline
mu & 8.4 & 5.4 & -2.4 & 8.2 & 18.8 & 1.0 \\\
theta[1] & 12.5 & 8.7 & -2.3 & 11.3 & 33.1 & 1.0 \\\
theta[2] & 8.1 & 6.7 & -5.0 & 8.0 & 21.7 & 1.0 \\\
theta[3] & 6.2 & 8.4 & -12.1 & 6.6 & 22.0 & 1.0 \\\
theta[4] & 7.9 & 6.8 & -5.9 & 8.1 & 21.6 & 1.0 \\\
theta[5] & 4.7 & 6.4 & -9.3 & 5.2 & 16.3 & 1.0 \\\
theta[6] & 6.1 & 7.1 & -9.3 & 6.4 & 19.5 & 1.0 \\\
theta[7] & 11.5 & 7.0 & -1.4 & 10.9 & 26.9 & 1.0 \\\
theta[8] & 8.9 & 8.3 & -8.0 & 8.7 & 27.3 & 1.0 \\\
tau & 7.9 & 5.8 & 1.4 & 6.5 & 22.0 & 1.0 \\\
\hline
\end{tabular}
\caption{Include a caption.}
\end{table}
```

Report Generation

- Just one line of code
- Report for N worst or all parameters of model
 - Or specific parameters
- On diagnostics, estimates or both
- PDF, Word or HTML
- ``generate_report()``

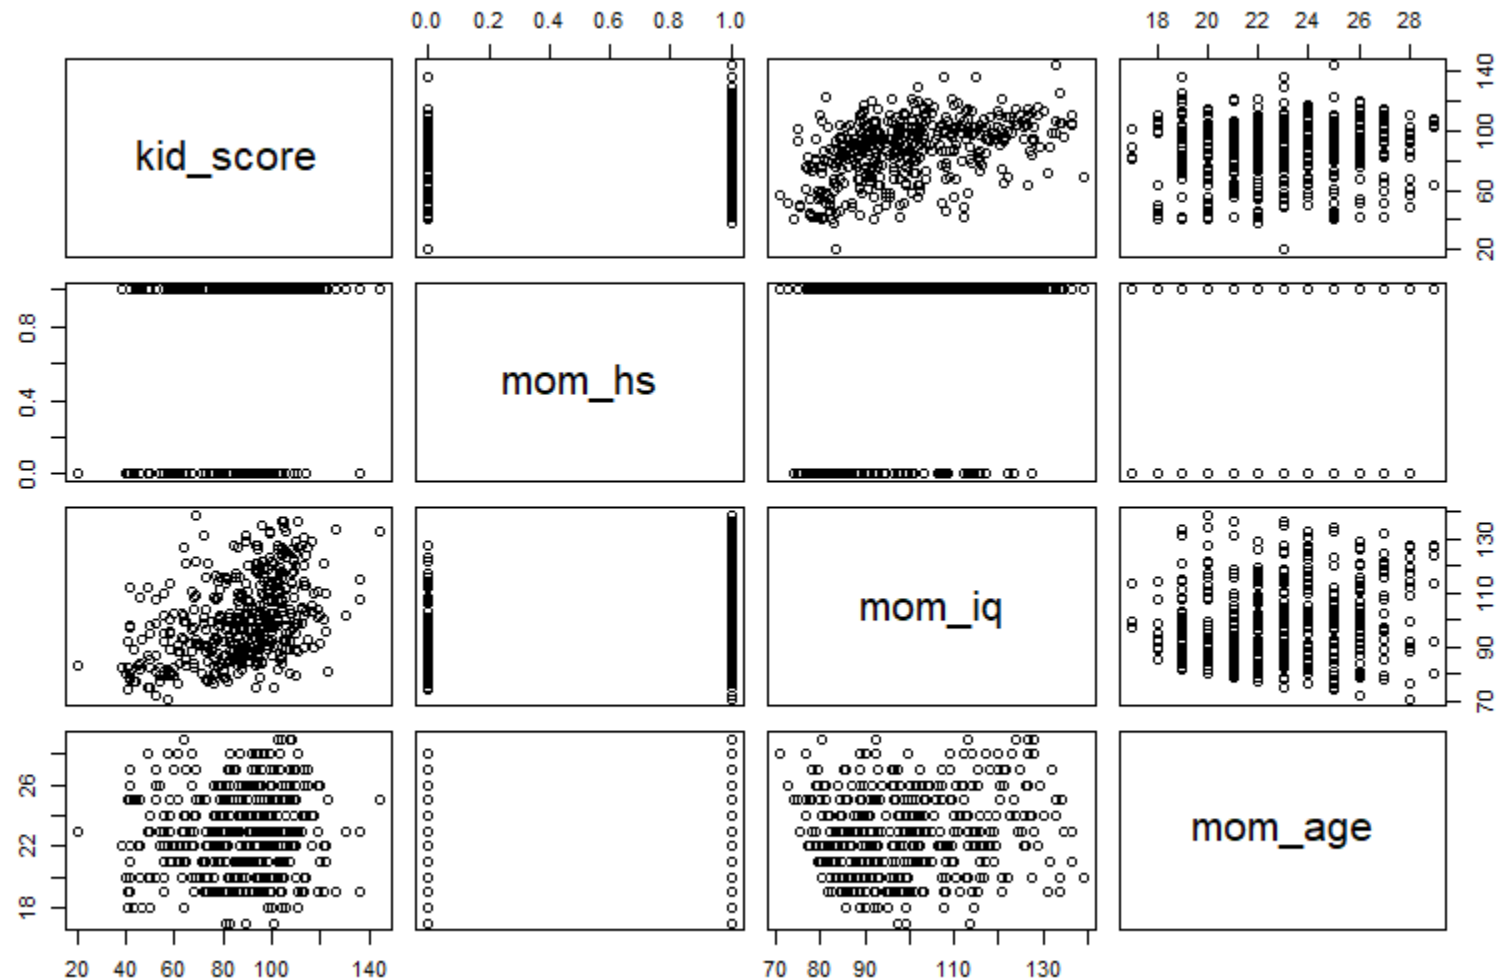
generate_report()

Arguments	
sso	A shinystan object.
n_param	On how many parameters do you want the report to be? To print the report for all parameters put 'n_param = Inf' or n_param = "all".
pars	An optional character vector of parameter names. If no names are specified the n_param count will be used. Arguments passed to pars take precedence over the n_param argument.
output_format	What type of report would you like? The options are 'html_document', 'pdf_document' and 'word_document'.
View	Do you want to open the report after it is generated?
report_type	What type of report would you like? The options are 'diagnose', 'estimate' and 'both'. The default is 'diagnose'.

Workflow

- Hopefully, live coding works
 - Or we have the next few slides
- We use a data set from rstanarm on IQ of kids.
 - `rstanarm::kidiq`

Visualize your data



Does it make sense?

$\text{IQ} \sim \text{Normal}(100, \text{sd} = 15)$

```
> kidiq$kid_score[order(kidiq$kid_score)]
[1] 20 38 40 41 41 41 42 42 42 42 42 42 42 43 43 44 45 46 46 46 47 49 49 49 50 50 50 50 50 52 52
[32] 52 52 52 54 54 54 56 56 56 56 56 56 56 57 57 58 58 58 58 58 58 58 58 59 60 60 60 61 61 61 61
[63] 62 62 63 63 64 64 64 64 64 64 65 65 65 65 65 66 67 67 67 67 68 68 68 68 68 69 69 69 69 69
[94] 70 70 71 72 72 72 72 73 73 73 73 73 73 74 74 74 74 75 75 75 75 76 76 76 76 76 76 76 76 77 77
[125] 77 77 78 78 78 78 78 78 78 78 79 79 79 79 79 80 80 80 80 80 81 81 81 81 81 82 82 82 82 83 83
[156] 83 83 83 83 83 83 83 83 84 84 84 84 84 84 84 85 85 85 85 85 85 86 86 86 86 86 86 86 86 86 86
[187] 87 87 87 87 87 87 87 87 87 87 87 87 87 87 88 88 88 88 88 88 89 89 89 89 89 89 89 90 90 90 90
[218] 90 90 90 90 91 91 91 91 91 92 92 92 92 92 92 92 92 92 92 93 93 93 93 93 94 94 94 94 94 94
[249] 94 94 94 94 94 94 94 94 94 94 94 95 95 95 95 95 95 95 95 95 95 96 96 96 96 96 96 97 97 97
[280] 97 97 97 97 97 97 98 98 98 98 98 98 98 98 98 98 99 99 99 99 99 99 99 99 99 99 100 100 100 100
[311] 100 100 100 100 100 100 100 101 101 101 101 101 101 101 102 102 102 102 102 102 102 102 102 102 102 102 102 102 103 103 103 103
[342] 104 104 104 104 104 104 104 104 104 104 104 104 104 104 104 105 105 105 105 105 106 106 106 106 106 106 106 106 106 107
[373] 107 107 107 107 107 108 108 108 108 109 109 109 109 109 109 109 110 110 110 110 110 110 110 111 111 111 111 112 112 112 112
[404] 113 113 113 113 113 113 114 114 114 115 115 116 116 116 117 117 119 119 119 120 120 121 121 122 122 123 126 130 136 136 144
> pnorm(20, 100, 15)
[1] 4.821303e-08
> pnorm(50, 100, 15)
[1] 0.0004290603
```

Does it make sense?

IQ \sim Normal(100, sd = 15)

```
> psych::describe(kidiq)
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
kid_score	1	434	86.80	20.41	90.00	87.93	19.27	20.00	144.00	124.00	-0.46	-0.19	0.98
mom_hs	2	434	0.79	0.41	1.00	0.86	0.00	0.00	1.00	1.00	-1.39	-0.07	0.02
mom_iq	3	434	100.00	15.00	97.92	99.11	15.89	71.04	138.89	67.86	0.47	-0.59	0.72
mom_age	4	434	22.79	2.70	23.00	22.71	2.97	17.00	29.00	12.00	0.18	-0.65	0.13

Decide on model

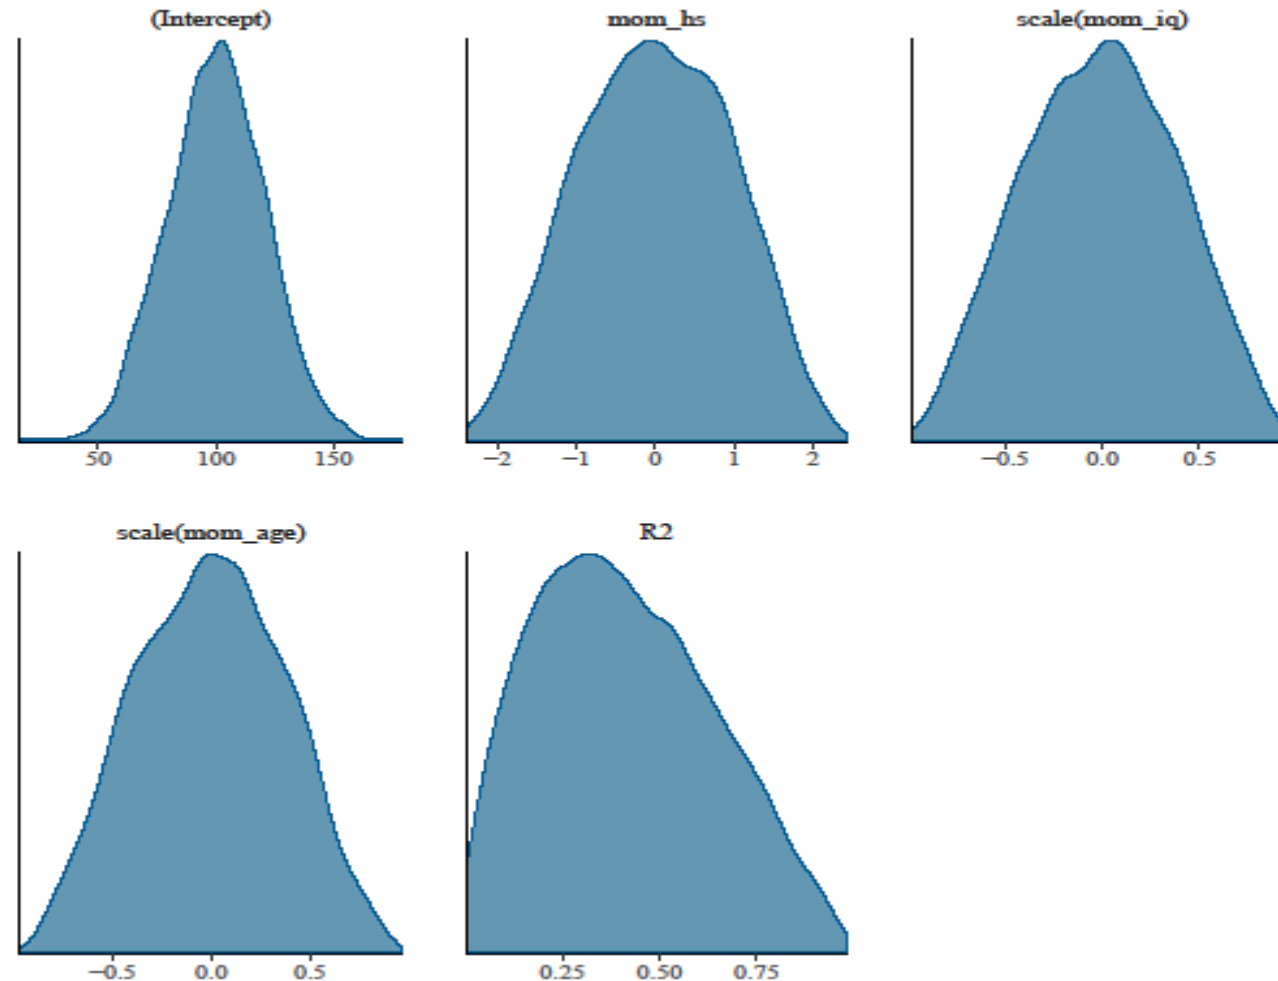
- Kidiq as outcome
- Simple regression
- Predictors
 - Did mom go to high school yes/no
 - IQ of mom – standardize?
 - Age of mom – standardize?
- We can use rstanarm!

Prior Predictive Distributions

uniform prior on R^2 if NULL, .3 R^2 seems reasonable in social sciences

```
fit_priorPD <- stan_lm(kid_score ~ mom_hs + mom_iq + mom_age,  
  data = kidiq,  
  prior = R2(.3),  
  prior_intercept = normal(100, 20),  
  prior_PD = TRUE,  
  seed = 13031990)
```


Prior Predictive Distributions



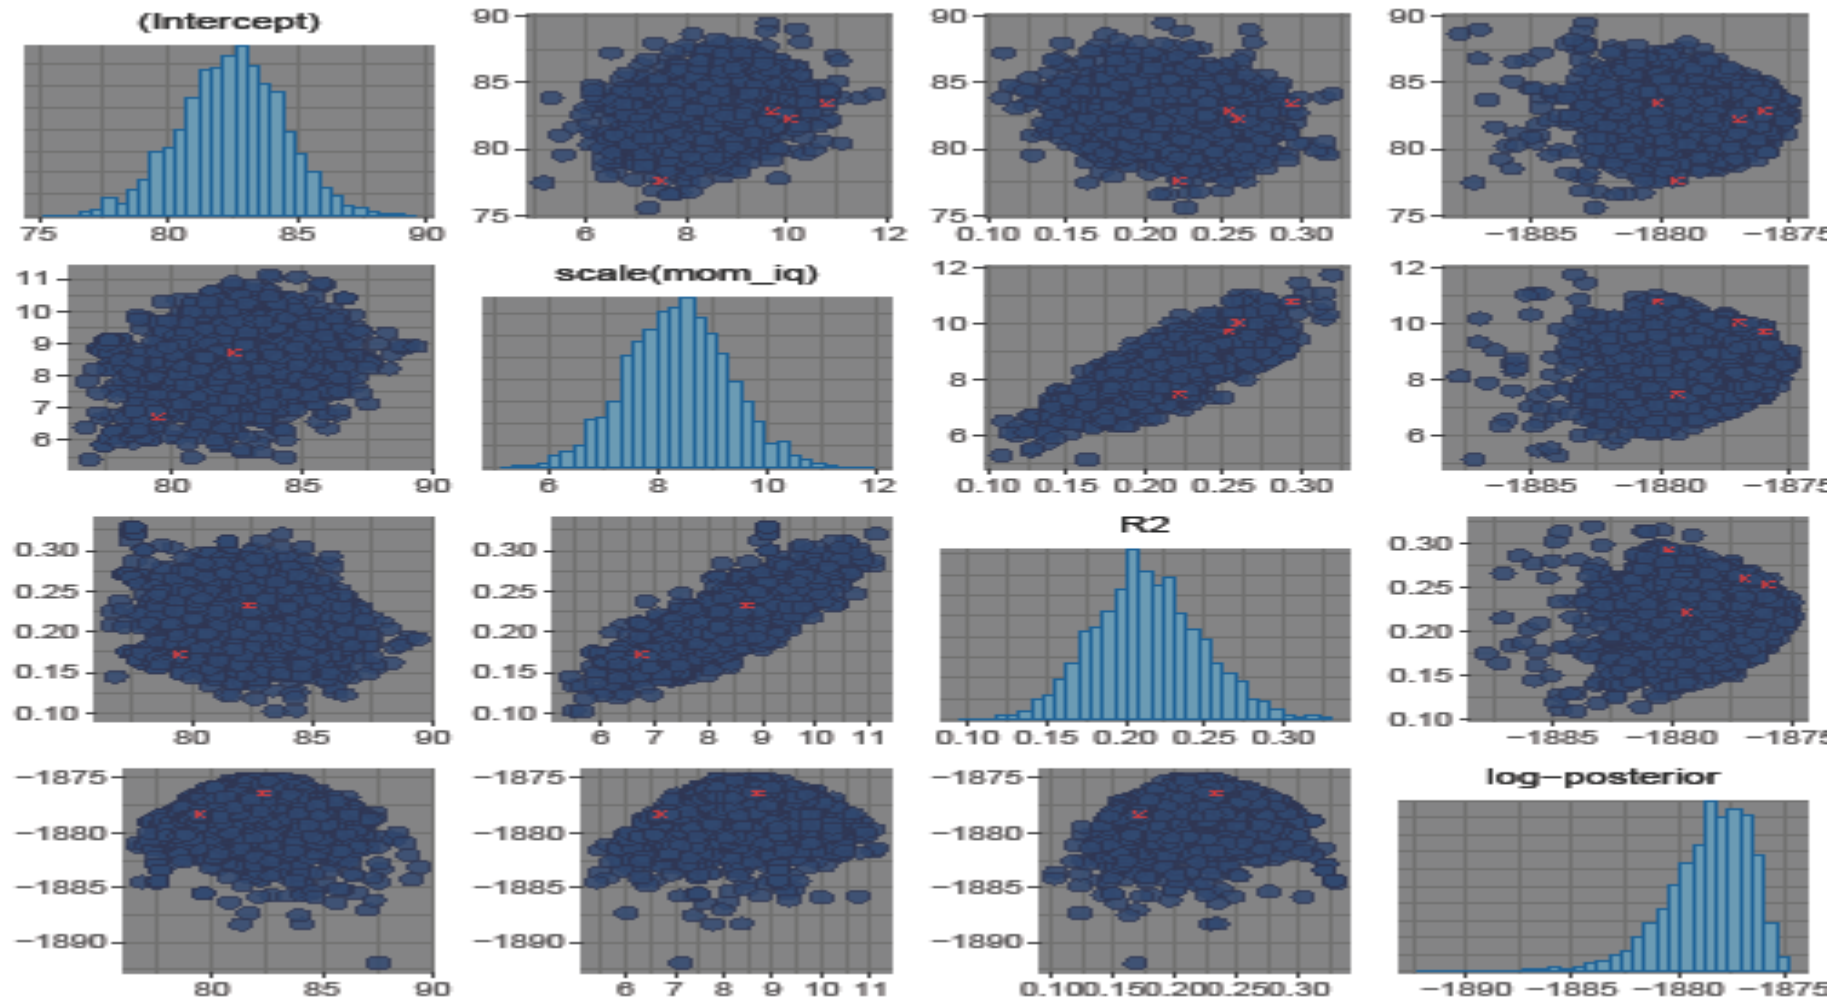
Estimate your model

```
fit <- stan_lm(kid_score ~ mom_hs + scale(mom_iq) + scale(mom_age),  
  data = kidiq,  
  prior = R2(.3), seed = 123456, iter = 2000, adapt_delta = .7,  
  prior_intercept = normal(100, 20),  
  prior_PD = FALSE)
```

Diagnose




Diagnose




Run again

```
fit <- stan_lm(kid_score ~ mom_hs + scale(mom_iq) + scale(mom_age),  
  data = kidiq,  
  prior = R2(.3), seed = 123456, iter = 2000, adapt_delta = .99,  
  prior_intercept = normal(100, 20),  
  prior_PD = FALSE)
```

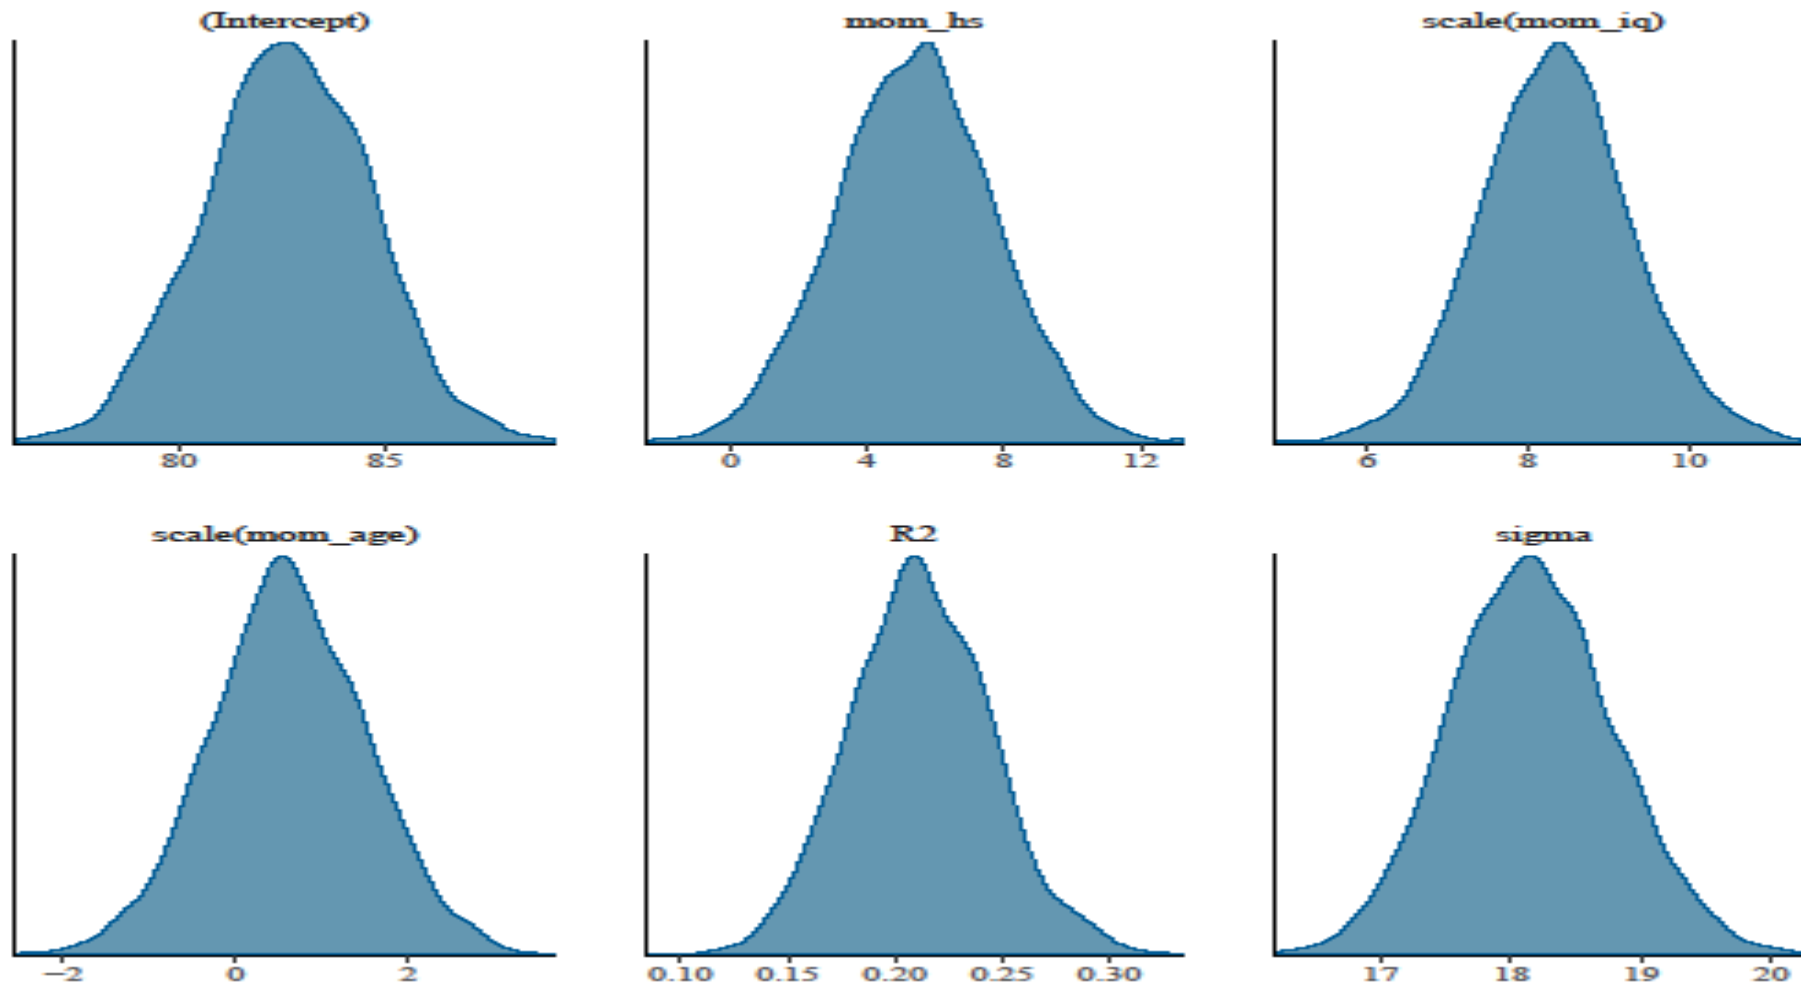
Diagnose again



Diagnose again




Look at estimates



Look at estimates

	mean	sd	2.5%	50%	97.5%	n_eff	Rhat
(Intercept)	82.62	1.99	78.73	82.63	86.41	2525.23	1.00
mom_hs	5.39	2.28	0.98	5.43	9.79	2554.19	1.00
scale(mom_iq)	8.34	0.92	6.57	8.34	10.23	2322.26	1.00
scale(mom_age)	0.61	0.90	-1.18	0.59	2.36	3330.14	1.00
sigma	18.16	0.62	16.98	18.15	19.41	3105.02	1.00

Look at estimates




Posterior Predictive

- Not in shinystan 3.0 at the moment
- Bayesplot very useful:
 - `posterior_predict()`
 - `ppc_.....`


```
fit_postPD <- posterior_predict(fit)
```

```
bayesplot::ppc_dens_overlay(kidiq$kid_score, fit_postPD, n_dens = 100)
```


Posterior Predictive




Posterior Predictive




Posterior Predictive



Posterior Predictive



Posterior Predictive



Wishlist

- Get you all to break it....

Try it out:

- `remove.packages("shinystan")`
- `devtools::install_github("stan-dev/shinystan@testing-modules", local = TRUE)`
- `shinystan::launch_shinystan_demo()`